

Unsupervised Learning -

Clustering

A clustering algorithm looks at a number of data points and automatically finds data points that are related or similar to each other.

Grouping unlabeled examples.

K-Means Algorithm

(centroid)

Step 0: Take 2 random points, to which can be the centre of the 2 clusters

Step 1: Assign each point to its closest centroid

Step 2: Recompute the centroid (find the average of all points that were assigned to it)

Go to Step 1

(If no further changes occur, stop algorithm)

Note: No. of centroids are k for k clusters.

Points are denoted as $x^{(1)}, x^{(2)}, x^{(3)}, \dots$

Randomly initialize K cluster centroids $\mu_1, \mu_2, \dots, \mu_K$

Repeat {

assign points to cluster centroids

for $i = 1 \text{ to } m$

$c^{(i)} :=$ index (from 1 to k) of cluster centroid closest to $x^{(i)}$

move cluster centroid.

for $k = 1 \text{ to } K$

$\mu_k =$ average (mean) of points assigned to cluster k

}

K-Means optimization objective.

$c^{(i)}$ → index of cluster ($1, 2, \dots, K$) to which example $x^{(i)}$ is currently assigned

μ_k = centroid of cluster centroid k

$\mu_{c(i)}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned

$$\begin{aligned} \text{Cost function} \Rightarrow J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) &= \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c(i)}\|^2 \\ &= \frac{1}{m} \sum_{i=1}^m \underbrace{\|x^{(i)} - \mu_{c(i)}\|^2}_{\text{Distortion function}} \end{aligned}$$

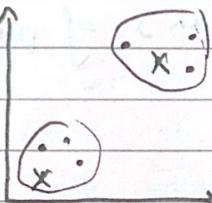
Initializing K-Means

↳ Random Initialization

Choose $K < m$

Randomly pick K training examples

Set μ_1, \dots, μ_K equal to these K examples



This method can create different clusters & all these sets of clusters may not be optimum.

Thus, we run it n times & find cost function (J) of all sets & choose the one with minimum J .

For $i = 1 \text{ to } 100$ &

Randomly initialize K-means.

Run K-means. Get $C^{(1)}, \dots, C^{(m)}, \mu_1, \mu_2, \dots, \mu_K$

Compute cost function (distortion)

$J(C^{(1)}, \dots, C^{(m)}, \mu_1, \mu_2, \dots, \mu_K)$

b

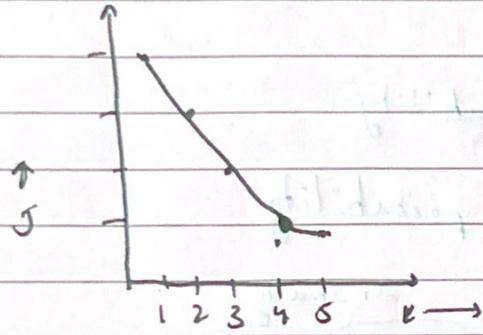
Pick set of clusters that gave lowest cost

Choosing Number of clusters

Elbow method \rightarrow drawing a graph

Elbow method \rightarrow drawing a graph against of J against K .

Where a sudden change occurs
we take that as K



$\therefore K = 4$
because J decreases
slowly after $K=4$

Best method to choose k
is to see what we
use the K value depending
on how we use the
clusters

(3 of 5 sizes for shirts)

$S, M, L \uparrow$
 XS, S, M, L, XL

common use

Fraud Detection
Manufacturing

Anomaly Detection

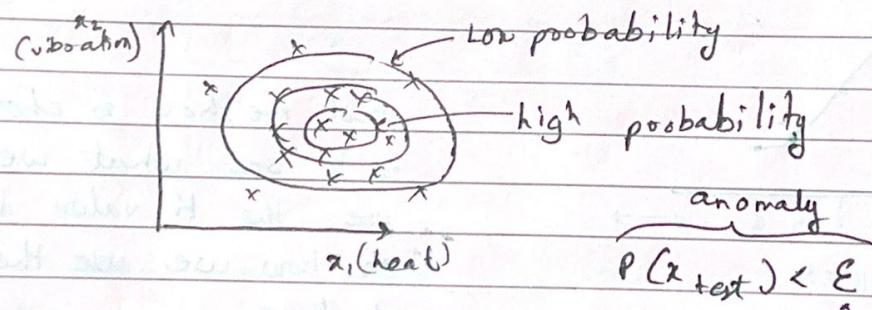
Basically when a new input of data differs from the usual pattern.

Density Estimation

Given a training set of m examples, we first train a model for probability of x being in dataset

Dataset: $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$

Model $p(x)$:

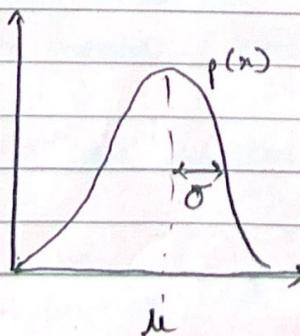


Gaussian (Normal) Distribution

Say x is a number

Probability of x is determined by a Gaussian with mean μ , variance σ^2

$$p(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$



Parameter Estimation for normal distribution

$$\mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

$$\sigma^2 = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu)^2$$

Anomaly Detection Algorithm

Training set: $\{\vec{x}^{(1)}, \vec{x}^{(2)}, \vec{x}^{(3)}, \dots, \vec{x}^{(m)}\}$

Each example $\vec{x}^{(i)}$ has n features

represented as x_n

$$\therefore p(\vec{x}) = p(x_1; \mu_1, \sigma_1^2) * p(x_2; \mu_2, \sigma_2^2) * p(x_3; \mu_3, \sigma_3^2)$$

$$= p(x_1; \mu_1, \sigma_1^2) * p(x_2; \mu_2, \sigma_2^2) * p(x_n; \mu_n, \sigma_n^2)$$

$$= \prod_{j=1}^n p(x_j; \mu_j, \sigma_j^2)$$

where

$$\mu_j = \frac{1}{m} \sum_{i=1}^m x_j^{(i)}$$

$$\sigma_j^2 = \frac{1}{m} \sum_{i=1}^m (x_j^{(i)} - \mu_j)^2$$

Evaluation in Anomaly Detection System.

Reinforcement Learning:

Type of ML that makes decisions by performing actions in an environment to achieve a goal.

The agent learns through trial & error.

Reward - correct action - +1

Penalty - negative action - -100

Basic flow

- Code starts at state 's'
- Takes an action 'a'
- Gets a reward 'R(s)'
- Reaches new state 's'

Return in RL

(If you have 5 dollars in front of you & 10\$ at a distance 30 min away, which one will you choose?)

Return essentially helps us decide which action has a better value depending on convenience

In simpler terms, return is a culmination of all rewards

↳ Discount factor is a small number (<1) which is added to the return statement which helps it prioritize quick gains

$$\therefore \text{Return} \Rightarrow R_1 + \gamma R_2 + \gamma^2 R_3 + \gamma^3 R_4 + \gamma^4 R_5 + \dots$$

γ = discount factor

Policy
set of rules that defines the actions an agent takes to maximize its return

if state s policy π action a

A policy is a function $\pi(s) = a$ mapping from state to actions - Tells us what action to take for given state s

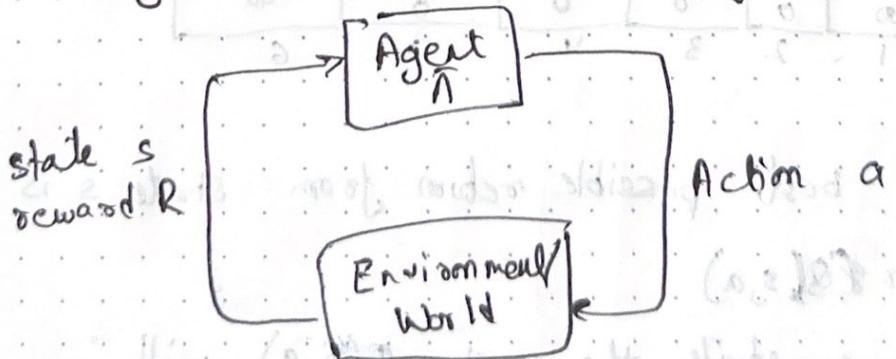
States \rightarrow representation of possible conditions agent is in.

Reward \rightarrow feedback from environment

Action \rightarrow Action is a decision to move

Markov Decision Process (MDP)
(formalization of all things based before)

(only current state matters for decision making)



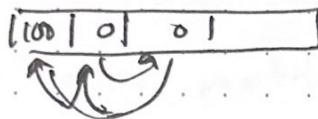
State Action value function (Q -function)

(allows us to calculate max. return for a state)

$Q(s, a) = \text{Return if you}$
• start in state s
• take action a (once)
• then behave optimally after
that

Ex $Q(2, \rightarrow) =$ go to new state k then
behave optimally to reach
desired target

State: 1 2 3



the $Q(2, \rightarrow) = 0 + (0.5)0 + (0.5)^2 0 + (0.5)^3 100$

= 12.5

Ex 2. $Q(2, \leftarrow) = 50$

So if we take an example

100	100	50	125	25	6.25	12.5	0	6.25	20	40	40
1	2	3	4	5	6	7	8	9	10	11	12

The best possible return from state s is

$$\max_a Q(s, a)$$

At a state the $\max Q(s, a)$ will

tell us the action ' a ' to take.

Q-function example

Bellman Equation

Helps compute the Q-Function (Breaks down Q-factors into a recursive problem)

We know that

$Q(s, a) = \text{Return if you start in state } s$
 • take action a once
 • then behave optimally after that

$s \rightarrow$ current state

$R(s) \rightarrow$ reward of current state

$a \rightarrow$ current action

$s' \rightarrow$ state you get to after action a

$a' \rightarrow$ action that you take in state s'



$$\text{Bellman Equation} \Rightarrow Q(s, a) = R(s) + (\gamma * \max_{a'} Q(s', a'))$$

Note: max
 a'
 max over all possible actions

Random (Stochastic) Environment

Continuous State Spaces

In previous example, we depicted distinct states.

but in actual application, these states are actually continuous.

Ex - a self driving car will be constantly on the move so the state will include its speed in X or Y axis, angle (θ), etc. position in X or Y axis, etc.

we can denote state as

$$\begin{bmatrix} x \\ y \\ \theta \\ v_x \\ v_y \end{bmatrix} \quad \text{speed}$$

State Value function for continuous state space

s → current state

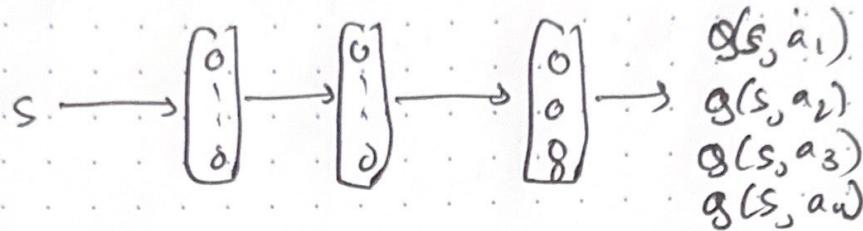
a → action

$$\vec{x} = \begin{bmatrix} s \\ a \end{bmatrix} \rightarrow \begin{bmatrix} 0 \\ \vdots \\ 0 \end{bmatrix} \rightarrow \begin{bmatrix} 0 \\ \vdots \\ 0 \end{bmatrix} \rightarrow \begin{bmatrix} 0 \end{bmatrix} \rightarrow Q(s, a) = y$$

We need to create a light NN to help compute $Q(s, a)$.

Notes if there are multiple actions, use binary encoding to specify action

A way to improve the Deep NN for more efficiency is to output all 'a' outputs.



So instead of taking inference of all possible actions in a state, we can simply pass the state & get output for all actions

ϵ -Greedy Policy

When the RL is still learning, how do we take action?

ϵ -Greedy Policy helps with this

In some state s

↳ with probability 0.95, pick the action a that maximises $q(s, a)$

↳ with probability 0.05, pick an action a randomly

$$\epsilon = 0.05$$

During training the NN may get stuck and not consider one option. So the 0.05 chance to choose a random action helps the code "explore" to find optimum solution